

ECE291

Lecture 7

Go set the table

Lecture outline

- Arrays
- Lookup tables
- Hash tables
- Jump tables
- Where this is leading...

Arrays

- Collection of sequentially addressable equally sized data in memory
- Only first element is labeled
- Index from there given size of each array element

Arrays - declaring

```
;declare and initialize a five element byte array  
MyByteArray DB    00h,0afh,12,83,83h
```

```
;declare a 50 element word array
```

```
MyWordArray1      RESW 50
```

```
MyWordArray2      RESB 100
```

```
MyWordArray3      RESD 25
```

```
;declare and initialize a 16 element char array
```

```
MyCharArray DB    "ECE291 Rocks!!!!", 0
```

Arrays - accessing

```
/* C array example */

int intWordArray[25];

for (int i = 0; i < 25; i++) {
    intWordArray[i] = i * i;
    /* Same as:
       *(intWordArray + 2*i) = i * i; */
}

/* This code segment declares a 25-element word
array (50 bytes) and fills each element with the
square of its index */
```

Arrays - accessing

;Assembly array example

```
intWordArray    RESW 25
```

```
mov cx, 25
```

```
begin_loop:
```

```
    mov bx, cx
```

```
    imul ax, cx, cx
```

```
    dec bx
```

```
    shl bx
```

```
    mov [intWordArray + bx], ax
```

```
loop begin_loop
```

Arrays in 2-dimensions

| | | | | | | | | | | | |
|----|----|----|----|----|----|----|----|----|----|----|----|
| 0 | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | A | B |
| C | D | E | F | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 |
| 18 | 19 | 1A | 1B | 1C | 1D | 1E | 1F | 20 | 21 | 22 | 23 |
| 24 | 25 | 26 | 27 | 28 | 29 | 2A | 2B | 2C | 2D | 2E | 2F |
| 30 | 31 | 32 | 33 | 34 | 35 | 36 | 37 | 38 | 39 | 3A | 3B |
| 3C | 3D | 3E | 3F | 40 | 41 | 42 | 43 | 44 | 45 | 46 | 47 |
| 48 | 49 | 4A | 4B | 4C | 4D | 4E | 4F | 50 | 51 | 52 | 53 |

- This is a 12 col by 7 row 2-D byte array
- It has $12 * 7 = 84 = 54h$ elements
- Addresses increase as you traverse a row
- Each row begins NUMCOLS addresses after the previous row
- Address of element (row, col) = row*NUMCOLS + col

```
NUMCOLUMNS EQU      12
NUMROWS EQU       7
My2DArray    RESB  NUMCOLUMNS * NUMROWS
```

```
mov bx, 4
mul bx, NUMCOLUMNS
add bx, 6

mov al, [My2DArray + bx]
; [My2DArray + 36] → AL
```

Lookup tables

- Lookup tables are special purpose arrays commonly used to convert one data form to another
- A lookup table is formed in the memory as a list of data that is referenced by a procedure to perform conversions

Lookup tables

- Consider movement in a 2-dimensional array by one element to an adjacent element (up, down, left, or right)

Array = Matrix of ROWSIZE x COLSIZE elements

Pos = Position in array = Row * COLSIZE + Column

Dir = Direction (UP=0, RIGHT=1, DOWN=2, LEFT=3)

- Slow and tedious original code compares each possible direction value

- 16 instructions

```
Mbegin
    CMP [DIR], UP
    JE MoveUP
    CMP [DIR], RIGHT
    JE MoveRT
    CMP [DIR], DN
    JE MoveDOWN
    CMP [DIR], LEFT
    JE MoveDown

MoveUP
    SUB [Pos], COLSIZE
    JMP MDone
MoveRT
    ADD [Pos], 1
    JMP MDone
MoveDN
    ADD [Pos], COLSIZE
    JMP MDone
MoveLT
    SUB [Pos], 1
    JMP MDone
MDone
```

Lookup tables

- Fast and Compact Table-lookup Code
- Eliminate conditional jumps by defining a table of movements indexed by the Direction
- 4 instructions + 4 word-sized variables

```
Movement
    dw -COLSIZE
    dw +1
    dw +COLSIZE
    dw -1
; Movement[0] == Go UP
; Movement[1] == Go RIGHT
; Movement[2] == Go DOWN
; Movement[3] == Go LEFT

Mbegin
    MOV BX, [DIR]           ; Direction = {0..3}
    SHL BX, 1               ; Index words, not bytes
    MOV AX, [Movement + BX]
    ADD [Pos], AX           ; Position += Movement
MDone: Finished!
```

Lookup tables

DAYS

```
PUSH DX  
PUSH SI  
  
MOV SI, DTAB    ;si points to table  
XOR AH, AH      ;clear AH  
ADD AX, AX      ;double AX  
ADD SI, AX      ;index to right day  
MOV DX, [SI]     ;get string address
```

```
MOV AH, 09h      ;display char magic  
INT 21h
```

```
RET
```

| | | |
|-----|----|---------------|
| SUN | DB | 'Sunday\$' |
| MON | DB | 'Monday\$' |
| TUE | DB | 'Tuesday\$' |
| WED | DB | 'Wednesday\$' |
| THU | DB | 'Thursday\$' |
| FRI | DB | 'Friday\$' |
| SAT | DB | 'Saturday\$' |

;Lookup table contains addresses of
;day of week strings. It converts
;0 to "Sunday", 1 to "Monday", etc.

```
DTAB DW SUN, MON, TUE,  
       WED, THU, FRI, SAT
```

Converts numbers 0 to 6
to days of the week
The number is passed
in AL register

Hash functions

- Consider using a lookup table when the input can span a large range of values but is sparsely populated.
- Example:
 - Count incoming packets from each of 100 hosts scattered throughout the Internet using the 32-bit IP source address
 - A Table-Lookup would require 4 billion table entries, most of which would be empty
- This is what we call “A Bad Thing.”

Hash functions

- Hash functions re-map large sparsely populated data sets (like the set of all IP addresses) into smaller more manageable data sets.

Large
sparsely
populated
data set

Hash Function

Small
tightly
packed
data set

Hash functions

- A Hash Function, $H(x)$, can be used to re-map the four-byte (W.X.Y.Z) IP address to a smaller (8-bit) value.
- $H(W.X.Y.Z) = W \text{ xor } X \text{ xor } Y \text{ xor } Z$

$$H(0.0.0.0) = 0$$

$$H(128.174.5.58) = 17$$

$$H(224.2.4.88) = 190$$

$$H(141.142.44.33) = 14$$

Hash Table

| Index | IP address |
|-------|---------------|
| 0 | 0.0.0.0 |
| ... | |
| 14 | 141.142.44.33 |
| ... | |
| 17 | 128.174.5.58 |
| ... | |
| 190 | 224.2.4.88 |
| ... | |

Hash functions

- Collisions can occur when two inputs map to the same output
 - H(128.174.112.1) and H(128.174.1.112) for example
- Collision resolution—Linear Probing
- Use next-available location in the hash table
- Add an extra column in the table for identifiers or tags to distinguish different entries in the table

```
tag = IP_addressTag  
index = h(IP_address);  
if (ht[index].tag == IP_addressTag) then  
    <Done>  
else  
    <search the hash table until the entry with  
    matching Tag is found>  
end if
```

Jump tables

- Suppose table entries are pointers to functions
- BX holds a command in the range of 0..N
- Let BX index a function

```
JFUNCTION          ; Input BX = Command
    SHL BX,1        ; Words, not bytes
    JMP [Jtable+BX] ; Jump to function
```

Jtable

| | |
|----|--------|
| dw | Funct0 |
| dw | Funct1 |
| .. | |
| dw | FunctN |

Funct0

; do something ..
RET

Funct1

; do something else ..
RET

FunctN

; do something else ..
RET

Where this is leading...

- Interrupts! These are triggers that cause the CPU to perform various tasks on demand
- Three kinds:
 - Software interrupts
 - Hardware interrupts
 - Exceptions
- Regardless of source, they are handled the same
 - Each interrupt has a unique interrupt number from 0 to 255. These are called interrupt vectors.
 - For each interrupt vector, there is an entry in the interrupt vector table.
 - The interrupt vector table is simply a jump table containing segment:offset addresses of procedures to handle each interrupt
 - These procedures are called *interrupt handlers* or *interrupt service routines*

Where this is leading...

- The first 1024 bytes of memory (addresses 00000 – 003FF) always contain the interrupt vector table. Always.
- Each of the 256 vectors requires four bytes—two for segment, two for offset

| Memory address (hex) | Interrupt function pointer |
|----------------------|----------------------------|
| 003FC | INT 255 |
| 4 * x | INT x |
| 00008 | INT 2 |
| 00004 | INT 1 |
| 00000 | INT 0 |

Software interrupts

- The operating system has handlers for many interrupts. These routines provide various built-in functions:
 - Displaying text to the screen (formerly known as “magic”)
 - Opening files
 - Rebooting the system (black magic)
 - Changing video modes
- Accessed with the INT instruction

DOS function dispatcher

- INT 21h is the DOS function dispatcher. It gives you access to dozens of functions built into the operating system.
- To execute one of the many DOS functions, you can specify a sub-function by loading a value into AH just before calling INT 21
- INT 21h sub-functions
 - AH=3Dh: Open File
 - AH=3Fh: Read File
 - AH=3Eh: Close File
 - AH=13h: Delete File (!)
 - AH=2Ah: Get system date
 - AH=2Ch: Get system time
 - AH=2Ch: Read DOS Version
 - AH=47h: Get Current Directory
 - AH=48h: Allocate Memory block (specified in paragraphs==16 bytes)
 - AH=49h: Free Memory block
 - AH=4Ch: Terminate program (and free resources)

System BIOS functions

- All PCs come with a BIOS ROM (or EPROM).
- The BIOS contains procedures that provide basic functions such as bootstrapping and primitive I/O.
 - INT 19h: Reboot system
 - INT 11h: Get equipment configuration
 - INT 16h: Keyboard I/O

Video BIOS functions

- Video cards come with procedures stored in a ROM
- Collectively known as the video BIOS
- Located at C0000-C7FFF and holds routines for handling basic video adapter functions
- To execute a function in video BIOS ROM, do an INT 10h with video sub-function number stored in AX
- INT 10h, Sub-function examples
 - AH=0, AL=2h: 80 column x 25 row text display mode
 - AH=0, AL=13h: 320x200 pixel, 256-color graphics display mode

Where is all this stuff???

