

Networking Lecture

Today's Topics

- [Computer Networking](#)
 - [Internet Protocols](#)
 - [IP Sockets Programming](#)
 - [NetBios Programming \(Obsolete, Broken on Windows 2000\)](#)
 - [Advanced Networks](#)
-

Computer Networks

- **Local-Area and Wide-Area Networks (LANs/WANs)**
Ethernet: 10/100/1000 Mbps (Local-Area)
 - Packets broadcast on shared media (the 'Ether')
 - Coax (10Base-2): Terminators prevent signal reflection
 - Twisted Pair (10/100/1000Base-T): Hubs or switches retransmit packets to links
 - No Quality of Service (QoS) Guarantees
- **Data Delivery**
 - Broadcast: One-to-all
 - Multicast: One-to-many (using group name)
 - Unicast: One-to-one (each host has unique name)
- **Data Transport Protocols**
 - Datagram
 - Packets may be lost
 - Packets may be reordered (except ATM)
 - Can Broadcast, Multicast, Unicast
 - Example: UDP
 - Connection-Oriented
 - Reliable connection between endpoints
 - Lost packets are retransmitted
 - Limited to point-to-point (Unicast) links
 - Example: TCP
- **Programming Models**
 - Client/Server (typical case)
 - Server is a centralized, permanent resource
 - Clients are short-lived, stateless programs
 - Unicast Connections between server and client
 - Peer-to-Peer

- All Programs share information
 - Multicast or broadcast data between hosts
 - **Network Adapter & Software**
 - Hardware address filters
 - Interrupt On Incoming Packet
 - Demultiplex incoming packet
 - Deliver data to appropriate application / higher level protocol (eg IP)
-

Internet Protocols

- **Internet Protocol (IP) is a global standard**
 - Allows for a common address space across entire Internet
 - 32-bit address, but not every address is usable
 - Packetized: data is split into segments for transmission
 - Unreliable transport, so packets may be:
 - lost
 - duplicated
 - arrive out of order
 - Independent of underlying hardware (may run on ethernet, ATM, etc)
 - We don't use IP directly, we mostly use two protocols built on it
 - **UDP (UDP/IP) - *User Datagram Protocol***
 - **Still unreliable**, but allows for multiplexing
 - Multiplexing = sharing of a single IP address by multiple programs
 - Performed by using a **Port number**, a 16 bit unsigned integer
 - Can have different source and destination ports
 - The operating system demultiplexes each incoming packet:
 - looks at the port number
 - sends the packet to the program listening on that port
 - Unidirectional, application needs to establish reverse path
 - Checksummed to prevent data corruption in packet
 - But is still unreliable (packets may be lost), just like pure IP
 - We'd like a reliable protocol for sending large amounts of data
 - **TCP (TCP/IP) - *Transmission Control Protocol***
 - A **reliable** stream transport protocol
 - Streamed connection: don't really see individual packets
 - Packets are guaranteed to arrive in the order sent, with no losses
 - Provides a "virtual circuit" connection
 - Full duplex once connection is established
 - Disadvantage: variable delay (sometimes *very* variable)
 - Has multiplexing, just like UDP (16-bit port numbers)
 - For how this actually works internally, take ECE 338!
-

IP "Sockets" Programming

- **So how do we actually *use* TCP/IP and UDP/IP?**
 - Most used: Sockets API
 - Introduced in 1981 in BSD 4.1 (Unix clone)
 - Implemented as system library calls
 - TCP and UDP interfaces very similar
 - Portable Microsoft implementation: WinSock
- **Byte Order**
 - As you all know, x86 machines are little endian.
 - IP protocols use network order, which is big endian.
 - Use functions to convert between "host" (little endian) byte order and network byte order:
 - [Socket_htonl](#), [Socket_htons](#) - convert 32-bit long, 16-bit short from host to network order
 - [Socket_ntohl](#), [Socket_ntohs](#) - vice-versa
 - **All socket functions require parameters (particularly the address and port) in network order**
- **[SOCKADDR](#) structure**
 - Contains both the port ([SOCKADDR.Port](#)) and the IP address ([SOCKADDR.Address](#))
 - Many sockets functions take a pointer to this structure instead of two separate values
- **Address Conversion**
 - Need a way to convert to and from human-readable (ASCII) names to IP addresses.
 - Some functions:
 - [Socket_inet_addr](#): convert a dotted decimal ("127.0.0.1") into an 32-bit address
 - [Socket_inet_ntoa](#): convert an 32-bit address into dotted decimal ASCII
 - [Socket_gethostbyname](#): converts an english host name ("www.uiuc.edu") into an IP address
 - [Socket_gethostbyaddr](#): converts an IP address into an english hostname
 - [Socket_gethostname](#): gets the local host's name (call [gethostbyname](#) to get the local IP)
- **Socket Creation and Setup**
 - [Socket_create](#): creates a socket of a particular [type](#) (called [socket\(\)](#) in standard implementations)
 - Waiting for connection (server) side:
 - [Socket_bind](#): binds a socket to a local IP address and port number
 - [Socket_listen](#): puts socket into a passive state (waiting for a connection)
 - [Socket_accept](#): accepts a new connection (TCP only)
 - Connecting (client) side:
 - [Socket_connect](#): connect to another (remote) socket (TCP only)
- **Sending and Receiving Data**
 - TCP way (after connection is established)
 - TCP is streaming, so we can view the connection as just a stream of bytes
 - [Socket_recv](#): receive data
 - [Socket_send](#): send data
 - UDP way
 - Each transmission is a new packet (unreliable, so it may be lost)
 - [Socket_recvfrom](#): gets a single incoming packet
 - [Socket_sendto](#): sends a single packet to another (remote) socket
- **Closing a connection**
 - [Socket_close](#): closes a socket completely (works for both TCP and UDP)
 - [Socket_shutdown](#): closes one direction of the connection (TCP only)
- **Examples**
 - Sockets programming (and network programming in general) is difficult.
 - Examples are **very** helpful!

- Because sockets are a standard interface, many examples are available on the web (basically all in C):
 - [WinSock application samples](#)
 - [Tutorial: Datagrams](#)
 - [Tutorial: Connections](#)
 - [CIS 307: Sockets](#)
 - PModelLib also includes a few example programs in assembly (in the V:\ece291\pmodelib\examples directory)
 - tcpweb: retrieves the ECE 291 webpage and prints it to the screen
 - tcpsrv: a simple TCP server
 - tcpcli: a simple TCP client (connects to tcpsrv)
 - udpsrv: a simple UDP server
 - udpcli: a simple UDP client (connects to udpsrv)
 - **Tips**
 - Pick your protocol (TCP or UDP) carefully!
 - TCP tends to be best for:
 - transaction-based communications
 - sending a large stream of data that must be received perfectly (eg a file)
 - UDP tends to be best for:
 - Anything realtime (game updates, sound files, etc)
 - Structure data to transmit over the network: the smaller the better
 - Use callback function (interrupt-driven network I/O) to not stop your main loop
 - PModelLib [Socket_SetCallback](#) and [Socket_AddCallback](#) allow you to do this in PModelLib.
 - You can trigger on different sockets for different [events](#), but *all* events get handled by *one* callback function.
 - **References**
 - As sockets is a standard interface to using TCP/IP and UDP/IP, there are some excellent references:
 - [Windows Sockets 1.1 Reference](#)
 - [CS/ECE 338 Class Notes](#)
 - [PModelLib documentation](#)
 - NOTES:
 - [PModelLib Sockets functions](#) are based on WinSock
 - Note: some PModelLib versions of WinSock functions have been simplified to make them easier to use in assembly.
-

NetBIOS (Obsolete, Broken on Windows 2000)

- **Local Area Network *service***
 - Independent of underlying hardware
 - Ethernet
 - Token Ring
 - FDDI
 - Independent of underlying Protocols
 - Standalone packets
 - Encapsulation in IP
 - Designed for Local Area Network
 - Assumes single local network segment
 - Implementation
 - DOS Packet Driver
 - Service provided by WinNT, Win95, WFWG

- **NetBIOS Provides**
 - Naming Service
 - Provides Unique names on Network
 - Global Namespace
 - Multiple types of communication
 - Datagram Broadcast
 - Datagram Multicast
 - Datagram Unicast
 - Reliable Connection
 - **Communication with NetBIOS**
 - Interrupt 5C
 - **Network Control Block (NCB)**
 - Basic NetBIOS Data Structure
 - Command Field: Determines what action to take
 - Inputs/Outputs/Pointers: All passed through NCB
 - **NetBIOS Response Methods**
 - Blocking Functions: Wait until event occurs
 - CallBack Functions: Register function to call when event occurs
Operates just like an interrupt
 - **NetBIOS Names**
 - 16-character NetBIOS identifiers
 - Local Names: Unique across the LAN (enforced by NetBIOS)
 - Group Name: Shared among multiple machines (Multicast)
 - **Examples**
 - **Broadcast/Multicast Datagrams**
 - Sender & Receiver ADD NAME
 - *Loop while running*
 - Sender SENDS Broadcast Datagram
 - Receiver RECEIVES Datagram
 - Sender & Receiver REMOVE NAME
 - **Unicast Connection**
 - Sender & Receiver ADD NAME
 - Receiver LISTENS
 - Sender CALLS (*Connection Established*)
 - *Loop while running*
 - Sender SENDS
 - Receiver RECEIVES
 - Sender or Receiver HANGUP (*Connection Terminated*)
 - Sender & Receiver REMOVE NAME
 - **Complete Documentation:** [CBIS NetBios Programmers Reference](#)
-

- **Programming with NetBIOS (Doesn't work under Windows 2000)**

- **Lockwood's NetLIB:** (*API to make your life easy!*)
 - [NetLIB readme file](#)
 - [NetLIB: ASM source code & examples](#)
- **Important Variables**
 - **grp_name:** Multicast Group name (*Default=ECE291NetLib\$\$\$\$*)
 - **my_name:** My unique network id (*Default=ECE291Player0\$\$\$*)
 - **TXBuffer:** Transmission buffer
Load with data that you want to transmit then call *SendPacket*.
 - **RXBuffer:** Receive Buffer
Filled with data from incoming network packet network just before your *netpost* function is called.
- **Procedures**
 - **NetINIT:** Call at start of your program
 - **netpost:** Callback function called whenever a datagram arrives.
 - Called with
 - BX = pointer to receive buffer
 - AX = length of data
 - Because this routine is called from an interrupt, it must:
 - Preserve all registers
 - Avoid DOS and LIB291 calls
 - **SendPacket:** Call with AX = Length of TXBuffer to transmit data
 - **NetRelease:** Call at the end of your program
- **PModeLib also has similiar functions**

Advanced Networks

- **ATM: Asynchronous Transfer Mode** (Local & Wide-Area)
 - Data transmitted as 53-byte cells
 - Messages transmitted using Adapation Layer (AAL)
 - Transmits data at rate of 155 Mbps (OC3), 622 Mbps (OC12), 2.4 Gbps (OC48), or 10 Gbps (OC192) *per host* and shared on virtual circuits.
 - Preserves *Quality of Service*, allowing integrated data (voice, video, IP) to carried over common network.
 - See the [iPOINT Testbed](#) for information about ATM research on this campus.

[Return to ECE291 Lecture Index](#)